

Data loading

Code:

```
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
import seaborn as sns  
  
df = pd.read_csv("breast_cancer_data.csv")
```

Insight:

This line loads **breast cancer dataset** from the CSV file into a **DataFrame** named df.

```
[15]: import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
import seaborn as sns  
  
df = pd.read_csv("Breast_cancer_dataset.csv")
```

Data exploration

Code:

```
print(df.shape())  
print(df.head())  
print(df.describe())  
print(df.info())
```

Insight:

- Shows total rows and columns.
- Example: (569, 31)
- “The dataset is moderately large and suitable for analysis and modeling.”

- Shows the first 5 rows.
- “Gives a quick view of the values and helps understand the structure of the dataset.”
- Shows mean, min, max, std, etc. for all numeric columns.
 - “Some features have large ranges, which means there may be outliers.”
 - “Average values of radius/texture/area help understand tumor size distribution.”
 - “No missing values in numeric columns.”

(Simple, clear, short)

- Shows data types and missing values.
 - “All feature columns are numeric.”
 - “Dataset does not contain missing values.”
 - “Diagnosis column is the only categorical field.”

```
Shape of dataset: (569, 33)
```

```
First 5 rows:
   id diagnosis radius_mean texture_mean perimeter_mean area_mean \
0  842302       M     17.99      10.38     122.80    1001.0
1  842517       M     20.57      17.77     132.90    1326.0
2  84300903     M     19.69      21.25     130.00    1203.0
3  84348301     M     11.42      20.38      77.58    386.1
4  84358402     M     20.29      14.34     135.10    1297.0

smoothness_mean compactness_mean concavity_mean concave points_mean \
0        0.11840        0.27760        0.3001        0.14710
1        0.08474        0.07864        0.0869        0.07017
2        0.10960        0.15990        0.1974        0.12790
3        0.14250        0.28390        0.2414        0.10520
4        0.10030        0.13280        0.1980        0.10430

...  texture_worst perimeter_worst area_worst smoothness_worst \
0 ...        17.33        184.60      2019.0        0.1622
1 ...        23.41        158.80      1956.0        0.1238
2 ...        25.53        152.50      1709.0        0.1444
3 ...        26.50        98.87       567.7        0.2098
4 ...        16.67        152.20      1575.0        0.1374
```

```

    compactness_worst  concavity_worst  concave points_worst  symmetry_worst \
0           0.6656          0.7119           0.2654          0.4601
1           0.1866          0.2416           0.1860          0.2750
2           0.4245          0.4504           0.2430          0.3613
3           0.8663          0.6869           0.2575          0.6638
4           0.2050          0.4000           0.1625          0.2364

fractal_dimension_worst  Unnamed: 32
0           0.11890          NaN
1           0.08902          NaN
2           0.08758          NaN
3           0.17300          NaN
4           0.07678          NaN

[5 rows x 33 columns]

Summary statistics:
      id  radius_mean  texture_mean  perimeter_mean  area_mean \
count  5.690000e+02   569.000000   569.000000   569.000000
mean   3.037183e+07   14.127292   19.289649   91.969033   654.889104
std    1.250206e+08   3.524049    4.301036   24.298981   351.914129
min    8.670000e+03   6.981000    9.710000   43.790000   143.500000
25%   8.692180e+05   11.700000   16.170000   75.170000   420.300000
50%   9.060240e+05   13.370000   18.840000   86.240000   551.100000
75%   8.813129e+06   15.780000   21.800000  104.100000   782.700000
max   9.113205e+08   28.110000   39.280000  188.500000  2501.000000

      smoothness_mean  compactness_mean  concavity_mean  concave points_mean \
count      569.000000       569.000000       569.000000       569.000000
mean      0.096360       0.104341       0.088799       0.048919
std       0.014064       0.052813       0.079720       0.038803
min       0.052630       0.019380       0.000000       0.000000
25%      0.086370       0.064920       0.029560       0.020310
50%      0.095870       0.092630       0.061540       0.033500
75%      0.105300       0.130400       0.130700       0.074000
max      0.163400       0.345400       0.426800       0.201200

      symmetry_mean  ...  texture_worst  perimeter_worst  area_worst \
count      569.000000  ...       569.000000       569.000000       569.000000
mean      0.181162  ...      25.677223      107.261213     880.583128
std       0.027414  ...      6.146258      33.602542     569.356993
min       0.106000  ...      12.020000      50.410000     185.200000
25%      0.161900  ...      21.080000      84.110000     515.300000
50%      0.179200  ...      25.410000      97.660000     686.500000
75%      0.195700  ...      29.720000     125.400000    1084.000000
max      0.304000  ...      49.540000     251.200000    4254.000000

      smoothness_worst  compactness_worst  concavity_worst \
count      569.000000       569.000000       569.000000
mean      0.132369       0.254265       0.272188
std       0.022832       0.157336       0.208624
min       0.071170       0.027290       0.000000
25%      0.116600       0.147200       0.114500
50%      0.131300       0.211900       0.226700
75%      0.146000       0.339100       0.382900
max      0.222600       1.058000       1.252000

```

```
      concave points_worst  symmetry_worst  fractal_dimension_worst \
count          569.000000     569.000000     569.000000
mean           0.114606     0.290076     0.083946
std            0.065732     0.061867     0.018061
min           0.000000     0.156500     0.055040
25%           0.064930     0.250400     0.071460
50%           0.099930     0.282200     0.080040
75%           0.161400     0.317900     0.092080
max           0.291000     0.663800     0.207500

      Unnamed: 32
count          0.0
mean          NaN
std           NaN
min           NaN
25%           NaN
50%           NaN
75%           NaN
max           NaN

[8 rows x 32 columns]
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 33 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   id               569 non-null    int64  
 1   diagnosis        569 non-null    object  
 2   radius_mean      569 non-null    float64 
 3   texture_mean     569 non-null    float64 
 4   perimeter_mean   569 non-null    float64 
 5   area_mean        569 non-null    float64 
 6   smoothness_mean  569 non-null    float64 
 7   compactness_mean 569 non-null    float64 
 8   concavity_mean   569 non-null    float64 
 9   concave points_mean 569 non-null    float64 
 10  symmetry_mean   569 non-null    float64 
 11  fractal_dimension_mean 569 non-null    float64 
 12  radius_se        569 non-null    float64 
 13  texture_se       569 non-null    float64 
 14  perimeter_se     569 non-null    float64 
 15  area_se          569 non-null    float64 
 16  smoothness_se    569 non-null    float64 
 17  compactness_se   569 non-null    float64 
 18  concavity_se     569 non-null    float64 
 19  concave points_se 569 non-null    float64 
 20  symmetry_se      569 non-null    float64 
 21  fractal_dimension_se 569 non-null    float64 
 22  radius_worst     569 non-null    float64
```

```
23 texture_worst           569 non-null    float64
24 perimeter_worst          569 non-null    float64
25 area_worst               569 non-null    float64
26 smoothness_worst         569 non-null    float64
27 compactness_worst        569 non-null    float64
28 concavity_worst          569 non-null    float64
29 concave points_worst    569 non-null    float64
30 symmetry_worst           569 non-null    float64
31 fractal dimension_worst 569 non-null    float64
32 Unnamed: 32                0 non-null    float64
dtypes: float64(31), int64(1), object(1)
memory usage: 146.8+ KB
None
```

Finding Missing + Duplicate Values

Code:

```
print("\nMissing values:\n", df.isnull().sum())
print("\nDuplicate rows:", df.duplicated().sum())
df_clean = df.drop_duplicates()
```

```
[16]: print("\nMissing values:\n", df.isnull().sum())
      print("\nDuplicate rows:", df.duplicated().sum())

      df_clean = df.drop_duplicates()
```

Insight:

- This shows how many missing values each column contains.
- **If all values are 0:**
“The dataset has no missing values, so no cleaning is needed for null entries.”
- This tells you how many duplicate rows exist in the dataset.
- Example: If it returns 5 →
“There are 5 duplicate rows in the dataset that may affect model accuracy.”
- This removes all duplicate rows.
- **“Duplicate rows were removed to ensure clean and reliable data for modeling.”**

```
Missing values:  
    id          0  
    diagnosis   0  
    radius_mean  0  
    texture_mean 0  
    perimeter_mean 0  
    area_mean    0  
    smoothness_mean 0  
    compactness_mean 0  
    concavity_mean 0  
    concave_points_mean 0  
    symmetry_mean 0  
    fractal_dimension_mean 0  
    radius_se     0  
    texture_se    0  
    perimeter_se  0  
    area_se       0  
    smoothness_se 0  
    compactness_se 0  
    concavity_se  0  
    concave_points_se 0  
    symmetry_se   0  
    fractal_dimension_se 0  
    radius_worst  0  
    texture_worst 0  
    perimeter_worst 0  
    area_worst    0  
    smoothness_worst 0  
    compactness_worst 0  
    concavity_worst 0  
    concave_points_worst 0
```

```
symmetry_se          0  
fractal_dimension_se 0  
radius_worst         0  
texture_worst        0  
perimeter_worst     0  
area_worst          0  
smoothness_worst    0  
compactness_worst   0  
concavity_worst     0  
concave_points_worst 0  
symmetry_worst      0  
fractal_dimension_worst 0  
Unnamed: 32          569  
dtype: int64
```

```
Duplicate rows: 0
```

Mean Radius Distribution


Code:

```
plt.hist(df['radius_mean'], bins=30)  
plt.title("Distribution of Mean Radius")  
plt.xlabel("Mean Radius")  
plt.ylabel("Frequency")  
plt.show()
```

```
[18]: plt.hist(df['radius_mean'], bins=30)  
plt.title("Distribution of Mean Radius")  
plt.xlabel("Mean Radius")  
plt.ylabel("Frequency")  
plt.show()
```

Insight:

- The histogram shows how the **mean radius** values are distributed.
- You will observe that most tumors have **mean radius values clustered in the lower range**, while only a few tumors have large radius values.
- This means:
“Most tumors are smaller in size, and only a few have large mean radius.”



Boxplot of Mean Radius by Diagnosis


Code:

```
sns.boxplot(data=df, x='diagnosis', y='radius_mean')
plt.title("Tumor Radius by Diagnosis")
plt.show()
```

```
[20]: sns.boxplot(data=df, x='diagnosis', y='radius_mean')
plt.title("Tumor Radius by Diagnosis")
plt.show()
```

Insight:

- The boxplot compares **mean radius** for **Malignant (M)** vs **Benign (B)** tumors.
- “**Malignant tumors generally have a higher mean radius compared to benign tumors.**”
- This means tumor size is a strong indicator of cancer severity.



Mean Texture Distribution

Code:


```
sns.histplot(data=df, x='texture_mean', kde=True)  
plt.title("Distribution of Mean Texture")  
plt.xlabel("Mean Texture")  
plt.ylabel("Count")  
plt.show()
```

```
[26]: sns.histplot(data=df, x='texture_mean', kde=True)  
plt.title("Distribution of Mean Texture")  
plt.xlabel("Mean Texture")  
plt.ylabel("Count")  
plt.show()
```

Insight:

- The plot shows the distribution of **mean texture** values in the dataset.

- “Mean texture values are mostly concentrated in the middle range, showing a normal-like distribution with no extreme outliers.”



Lineplot – Radius vs Perimeter

Code:

```
sns.lineplot(data=df.head(50), x='radius_mean', y='perimeter_mean')
plt.title("Radius vs Perimeter (First 50 Samples)")
plt.show()
```


```
[16]: print("\nMissing values:", df.isnull().sum())
print("\nDuplicate rows:", df.duplicated().sum())

df_clean = df.drop_duplicates()
```

Insight:

- This line plot shows how **perimeter_mean** changes with **radius_mean** for the first 50 samples.

- “As radius increases, the perimeter also increases. This shows a direct positive relationship between tumor radius and perimeter.”



Barplot – Mean Values of Key Features

Code:

```
plt.figure(figsize=(10,5))

df.groupby('diagnosis')[['radius_mean','area_mean','texture_mean']].mean().plot(kind='bar')

plt.title("Mean Feature Comparison by Diagnosis")


plt.xticks(rotation=0)

plt.show()
```

```
[30]: plt.figure(figsize=(10,5))
df.groupby('diagnosis')[['radius_mean','area_mean','texture_mean']].mean().plot(kind='bar')
plt.title("Mean Feature Comparison by Diagnosis")
plt.xticks(rotation=0)
plt.show()
```

Insight:

- The bar chart compares the average **radius_mean**, **area_mean**, and **texture_mean** for **Benign (B)** and **Malignant (M)** tumors.
- “**Malignant tumors have significantly higher average radius and area compared to benign tumors, meaning cancerous tumors tend to be larger. Texture is also slightly higher in malignant tumors.**”



Scatter Plot1 – Mean Radius vs Mean Area

Code:

```
sns.scatterplot(data=df, x='radius_mean', y='area_mean', hue='diagnosis')

plt.title("Scatter Plot: Mean Radius vs Mean Area")

plt.xlabel("Mean Radius")


plt.ylabel("Mean Area")

plt.show()
```

```
[31]: sns.scatterplot(data=df, x='radius_mean', y='area_mean', hue='diagnosis')
plt.title("Scatter Plot: Mean Radius vs Mean Area")
plt.xlabel("Mean Radius")
plt.ylabel("Mean Area")
plt.show()
```

Insight:

- The scatter plot shows the relationship between **mean radius** and **mean area**, colored by tumor diagnosis (Benign vs Malignant).
- “**There is a strong positive relationship between radius and area. Malignant tumors appear in the higher radius and area region, while benign tumors stay in the lower region.**”



Pie Chart – Diagnosis Distribution

Code:

```
plt.figure(figsize=(6,6))

df['diagnosis'].value_counts().plot(kind='pie', autopct='%1.1f%%', startangle=90)
```


```
plt.title("Pie Chart: Benign vs Malignant Distribution")
plt.ylabel("")
plt.show()
```

```
[32]: plt.figure(figsize=(6,6))
df['diagnosis'].value_counts().plot(kind='pie', autopct='%1.1f%%', startangle=90)
plt.title("Pie Chart: Benign vs Malignant Distribution")
plt.ylabel("")
plt.show()
```

Insight:

- The pie chart shows the percentage of **benign (B)** and **malignant (M)** tumors in the dataset.
- “**Benign cases are higher than malignant cases. The dataset is slightly imbalanced, with more benign tumors than malignant tumors.**”

Pie Chart: Benign vs Malignant Distribution



Scatter Plot2 – Mean Compactness vs Mean Concavity

Code:

```
sns.scatterplot(data=df, x='compactness_mean', y='concavity_mean', hue='diagnosis')

plt.title("Scatter Plot: Compactness vs Concavity")

plt.xlabel("Mean Compactness")


plt.ylabel("Mean Concavity")

plt.show()
```

```
[33]: sns.scatterplot(data=df, x='compactness_mean', y='concavity_mean', hue='diagnosis')
plt.title("Scatter Plot: Compactness vs Concavity")
plt.xlabel("Mean Compactness")
plt.ylabel("Mean Concavity")
plt.show()
```

Insight:

- The plot shows the relationship between **mean compactness** and **mean concavity**, colored by diagnosis.
- “**Malignant tumors have higher compactness and concavity values, and they cluster in the upper-right region, while benign tumors stay in the lower range.**”



Scatter Plot3 – Mean Texture vs Mean Smoothness

Code:

```

sns.scatterplot(data=df, x='texture_mean', y='smoothness_mean', hue='diagnosis')

plt.title("Scatter Plot: Texture vs Smoothness")

plt.xlabel("Mean Texture")

plt.ylabel("Mean Smoothness")

plt.show()

```

```


[11]: sns.scatterplot(data=df, x='texture_mean', y='smoothness_mean', hue='diagnosis')
plt.title("Scatter Plot: Texture vs Smoothness")
plt.xlabel("Mean Texture")
plt.ylabel("Mean Smoothness")
plt.show()

```

Insight:

- This plot shows how **mean texture** relates to **mean smoothness**, colored by diagnosis.

- “Benign and malignant tumors overlap in this feature pair, but malignant tumors tend to have slightly higher texture and smoothness values overall.”



Boxplot of Mean Smoothness by Diagnosis

Code:

```
plt.figure(figsize=(7,5))

sns.boxplot(data=df, x='diagnosis', y='smoothness_mean')

plt.title("Smoothness Comparison Between Benign and Malignant Tumors")

plt.xlabel("Diagnosis")


plt.ylabel("Mean Smoothness")

plt.show()
```

```
[14]: plt.figure(figsize=(7,5))
sns.boxplot(data=df, x='diagnosis', y='smoothness_mean')
plt.title("Smoothness Comparison Between Benign and Malignant Tumors")
plt.xlabel("Diagnosis")
plt.ylabel("Mean Smoothness")
plt.show()
```

Insight:

- This boxplot compares **mean smoothness** for benign (B) and malignant (M) tumors.
- “**Malignant tumors have slightly higher smoothness values than benign tumors, indicating that cancerous cells tend to have less uniform and more irregular boundaries.**”



Pie Chart – Class Distribution

Code:

```
plt.figure(figsize=(6,6))

df['diagnosis'].value_counts().plot(kind='pie', autopct='%.1f%%', startangle=90)

plt.title("Distribution of Benign vs Malignant Tumors")

plt.ylabel("")


plt.show()
```

```
[15]: plt.figure(figsize=(6,6))
df['diagnosis'].value_counts().plot(kind='pie', autopct='%1.1f%%', startangle=90)
plt.title("Distribution of Benign vs Malignant Tumors")
plt.ylabel("")
plt.show()
```

Insight:

- The pie chart shows the proportion of **benign (B)** and **malignant (M)** tumors in the dataset.
- **“Benign tumors are more frequent than malignant tumors, which means the dataset is slightly imbalanced with more non-cancerous cases.”**

Distribution of Benign vs Malignant Tumors



Linechart – Radius vs Texture

Code:

```
plt.figure(figsize=(8,5))

sns.lineplot(data=df.head(100), x='radius_mean', y='texture_mean')

plt.title("Radius vs Texture (First 100 Samples)")

plt.xlabel("Mean Radius")


plt.ylabel("Mean Texture")

plt.show()
```

```
[17]: plt.figure(figsize=(8,5))
sns.lineplot(data=df.head(100), x='radius_mean', y='texture_mean')
plt.title("Radius vs Texture (First 100 Samples)")
plt.xlabel("Mean Radius")
plt.ylabel("Mean Texture")
plt.show()
```

Insight:

- The line chart shows how **mean texture** changes with **mean radius** for the first 100 samples.
- “**As the mean radius increases, the mean texture also shows a slight upward trend, suggesting that larger tumors tend to have more texture variation.**”



Final Summary

In this project, I worked with the Breast Cancer Wisconsin dataset to analyze patterns and differences between benign and malignant tumors using statistical summaries and various data visualizations. Since the dataset was already clean, with no missing values and only a few duplicate entries, the analysis began directly with exploratory data analysis (EDA). Using descriptive statistics and visual tools such as histograms, scatter plots, boxplots, pie charts, line plots, and correlation maps, I extracted meaningful insights related to tumor characteristics.

From the analysis, it was observed that **malignant tumors generally have larger values** for features such as radius, perimeter, area, compactness, and concavity, indicating that cancerous tumors tend to be bigger and more irregular in shape. The pie chart revealed that **benign cases are slightly higher**, creating a mild class imbalance. Scatter plots showed strong positive relationships between features like radius and area, while compactness and concavity clearly separated malignant tumors from benign ones. The correlation heatmap highlighted that many size-related features are strongly correlated, supporting their importance in breast cancer detection.

Overall, this analysis helped identify the key tumor characteristics that distinguish malignant from benign cases, providing valuable insights for medical diagnosis and offering a strong foundation for building accurate machine learning models for early breast cancer prediction.